

Overview

My solution is built as a standalone **Next.js** application, utilizing **Azure OpenAI** for processing. The architecture was designed in order to meet the latency requirements of the field agents (5 seconds) and the strict network security constraints of a federal Azure environment.

Technical Stack

- **Framework:** Next.js (App Router)
- **AI Engine:** Azure OpenAI (GPT-4.1-mini)
- **Image Processing:** Sharp (for pre-processing/refinement)
- **Deployment Target:** Azure Static Web Apps (to ensure FedRAMP/Firewall compatibility)

Tools Used

- **Github Copilot**

Design Decisions

1. Sarah Chen's 5 Second Latency and Batch Upload Requirements.

To meet Sarah Chen's requirement for results in ~5 seconds, I opted to use GPT-4.1-mini, a lightweight model with high OCR benchmark scores, for the label extraction and validation process.

- **Why not a different model:** While "Reasoning" models (like GPT-5-Mini) offer higher accuracy, their 15s+ latency would lead to agents abandoning the solution.
- **Making GPT 4.1-mini more accurate:** I used a Parallel Ensemble Approach. By running three extraction calls simultaneously, and taking the best parts of each, I achieve the speed of a "mini" model with the accuracy of a much larger one through redundancy.
- **Batch Upload:** The system supports batch processing, allowing the upload of an image folder and a corresponding CSV file (ideally provided by the COLA system) containing the necessary application data. Labels are matched to the CSV data and processed in parallel batches of 5 to manage current API rate limits.

2. Marcus William's Firewall Requirement:

Following Marcus Williams' feedback regarding blocked outbound traffic, I chose not to use APIs outside of **Azure-native services**. This ensures that the prototype can be white-listed within the existing TTB infrastructure without the "scanning vendor disaster" of 2023.

3. Dave Morrison's "Human Judgement" and Jenny Park's exact Government Warning

To address Dave's concerns about nuance (e.g., "Stone's Throw" vs. "STONE'S THROW"), I implemented a dual-layer validation:

- **AI Judgment Layer:** For fields like Brand Name and Class/Type, the AI evaluates "semantic equivalence" rather than strict string matching.
- **Hard-Coded String Matching:** For the **Government Warning**, the system enforces 100% literal accuracy (caps/bolding/text), meeting Jenny Park's strict compliance standards.

The Extraction Pipeline: Parallel Consensus

Rather than relying on a single (potentially hallucinatory) pass, the system follows a three-step synthesis:

1. **Extraction (Parallel):** Three separate instances of the model extract data from the label. Each is given the "definition" of the fields (e.g., "What is a Brand Name?") without being told the "expected" answer to prevent confirmation bias/hallucination.
2. **Evaluation (Parallel):** A separate AI evaluates each extracted field against the application data, assigning a binary (0 or 1) score based on accuracy and compliance rules.
3. **Heuristic Synthesis:** A final layer selects the highest-scoring extraction for each category, combines them together then presents the final decision of if the label is valid or not to the user.

Assumptions, Limitations & Errors

Limitations:

- **API Rate Limits:** Current API request limits (100 requests per minute on the personal Azure account) restrict the speed of batch processing. While the system could potentially handle 100+ labels simultaneously with higher limits, it currently processes large batches in a parallel queue of 5 labels to manage the constraint.
- **Limited AI Model Options:** Azure AI API gives access to a limited number of models (my account only had access to OpenAI models). I would likely have used Gemini 3 Flash if I had access to the google library.

Assumptions:

- **COLA System CSV Export:** It is assumed that the COLA system can export a CSV file containing the expected label information. If this export functionality is unavailable, users will be required to manually input the necessary details for each bottle for batch requests.
- **Image Format:** Labels are assumed to be uploaded as a single continuous image, consistent with the examples provided on the TTB website.

Errors:

- **Azure API Content Filter:** Occasionally, the Azure AI filter may flag alcohol bottle labels as inappropriate, resulting in a 400 "Image processing blocked due to content policy violation" error. In the event this occurs, my system is designed to re-upload the image and retry the process, which typically resolves the issue.
- **Api Limit:** When the GPT-4.1-Mini hits an api call limit it throws a 500 error. My system starts to use a backup model, GPT 4.1, when this happens.

Example Walkthrough: “OLD TOM DISTILLERY BOTTLE”

Extracting OLD TOM DISTILLERY bottle label

Expected Details

- Brand Name: "OLD TOM DISTILLERY"
- Class/Type: "Kentucky Straight Bourbon Whiskey"
- Alcohol Content: "45% Alc./Vol. (90 Proof)"
- Net Contents: "750 mL"
- Government Warning: [Standard government warning text]

Step 1. Perform 3 extraction in parallel of uploaded label using GPT 4.1 Mini

Extraction 1	Extraction 2	Extraction 3
Brand Name: Old Tom Distillery	Brand Name: Distillery	Brand Name: idTom Dry
Class/Type : “Straight”	Class/Type: "Kentucky Straight Bourbon Whiskey"	Class/Type: “Kentucky”
Alcohol Content: “30% Alc./Vol. (60 Proof)”	Alcohol Content: “40% Alc./Vol. (80 Proof)”	Alcohol Content: “45% Alc./Vol. (90 Proof)”
Net Contents: "750 mL"	Net Contents: "300 mL"	Net Contents: "254 mL"

Government Warning : Issue with label capitalization	Government Warning : Perfect Label	Government Warning : Issue with label wording

Step 2. Evaluating each extraction against expected data in parallel (0 invalid label | 1 valid)

Evaluation 1	Evaluation 2	Evaluation 3
Brand Name: Old Tom Distillery 1	Brand Name: Distillery 0	Brand Name: IdTom Dry 0
Class/Type : "Straight" 0	Class/Type: "Kentucky Straight bourbon whiskey" 1	Class/Type: "Kentucky" 0
Alcohol Content: "30% Alc./Vol. (60 Proof)" 0	Alcohol Content: "40% Alc./Vol. (80 Proof)" 0	Alcohol Content: "45% Alc./Vol. (90 Proof)" 1
Net Contents: "750 mL" 1	Net Contents: "300 mL" 0	Net Contents: "254 mL" 0
Government Warning : Issue with label capitalization 0	Government Warning : Perfect Label 1	Government Warning : Issue with label wording 0

Step 3. Synthesize the best elements from each extraction and output that to the user

Combined Extraction : VALID
Brand Name: Old Tom Distillery 1
Class/Type: "Kentucky Straight Bourbon Whiskey" 1
Alcohol Content: "45% Alc./Vol. (90 Proof)" 1`
Net Contents: "750 mL" 1
Government Warning : Perfect Label 1